

## Document Search Service

A secure, multi-tenant document indexing and search service built with **Spring Boot**, **PostgreSQL**, **Elasticsearch**, and **Redis**.

The service supports **JWT-based authentication**, **role-based access control**, **tenant isolation**, and **high-performance search**.

## Features

- JWT authentication with role-based authorization (USER, ADMIN)
- Multi-tenant architecture with strict data isolation
- Document CRUD operations
- Full-text search powered by Elasticsearch
- Redis caching for faster search responses
- Tenant-level rate limiting
- Structured logging and observability-ready design

## Tech Stack

Layer	Technology
Backend	Java 17, Spring Boot
Security	Spring Security, JWT
Database	PostgreSQL
Search Engine	Elasticsearch
Cache	Redis
Build Tool	Maven
Containerization	Docker

## Architecture Overview

Client

|

| JWT



## Spring Boot Application

```
|— Security & RBAC  
|— Tenant Context  
|— Rate Limiting  
|— Cache Layer  
  
|  
|— PostgreSQL (Source of Truth)  
|— Elasticsearch (Search Index)  
└— Redis (Caching)
```

## API Endpoints

### Authentication

POST /auth/token

### Document APIs

POST /documents

GET /documents/{id}

DELETE /documents/{id} (ADMIN only)

### Search API

GET /search?q={query}&tenant={tenantId}

## Sample Requests

### Get JWT Token

POST /auth/token

Content-Type: application/json

{

  "username": "user1",

```
    "password": "password"  
}
```

## Create Document

POST /documents

Authorization: Bearer <JWT>

Content-Type: application/json

```
{  
    "title": "Sample Document",  
    "content": "This is a searchable document",  
    "tags": ["spring", "search"]  
}
```

## Search Documents

GET /search?q=spring&tenant=tenant1

Authorization: Bearer <JWT>

## Multi-Tenancy Model

- Tenant ID is embedded in JWT claims
- Extracted and enforced using TenantContext
- All database and search queries are tenant-scoped
- Cross-tenant access is explicitly blocked

## Caching Strategy

- Redis used as a read-through cache
- Cached search results with configurable TTL
- Cache invalidation on document create/delete
- Jackson JavaTimeModule enabled for date serialization

## **Security Model**

- Stateless JWT authentication
- Role-based access control at:
  - Route level
  - Method level (@PreAuthorize)
- Enforced tenant isolation
- HTTPS-ready configuration

## **Consistency Guarantees**

- Strong consistency in PostgreSQL
- Eventual consistency between PostgreSQL and Elasticsearch
- Cache TTL ensures freshness while improving performance

## **Running Locally**

### **Prerequisites**

- Java 17+
- Docker & Docker Compose
- Maven

### **Start Dependencies**

```
docker-compose up -d
```

### **Run Application**

```
mvn spring-boot:run
```

## **Production Readiness (Design Considerations)**

- Horizontal scaling with stateless services
- Elasticsearch index sharding
- Redis cluster support
- Retry and circuit breaker readiness

- Centralized logging & metrics integration
- Zero-downtime deployments

## Project Structure

src/main/java

```
|—— config  
|—— controller  
|—— security  
|—— service  
|—— repository  
|—— model  
└—— component
```

## Future Enhancements

- Async indexing with Kafka
- Advanced relevance tuning
- Pagination & sorting support
- Kubernetes deployment manifests
- OpenAPI (Swagger) documentation

## Author

**Aakanksha Saxena**

Senior Java Backend Engineer